

Design Justification Report

1. Problem and Motivation

The core objective of this project is the hardware acceleration of an indoor localization system (ILS) utilizing a Convolutional Neural Network (CNN) based spatial processing engine. The system architecture is fundamentally constrained by the specific data availability inherent to the target physical environment. During initial project scoping, it was determined that collecting Round-Trip Time (RTT) data was not feasible for the current setup. Consequently, the localization methodology must rely exclusively on Received Signal Strength Indicator (RSSI) data. Processing these 2D RSSI spatial maps to deduce physical positioning requires an extensive, continuous matrix of parallel dot-product calculations, forming the core mathematical kernel that our architecture must execute.

Executing this spatial processing engine strictly in software proved to be highly inefficient and unsuitable for edge deployment. The Milestone 1 software baseline benchmark was conducted utilizing a general-purpose processor, specifically an Intel Core i7-12700H CPU operating at a peak frequency of 4.7 GHz. The Python-based baseline model yielded a median execution time of 1.0918 seconds per epoch over ten consecutive runs. Translating this execution time to our primary performance metric resulted in a throughput of merely 183.18 RSSI samples processed per second. Furthermore, this software execution consumed approximately 150 MB of peak memory (Resident Set Size).

This performance profile is entirely inadequate for real-time localization tracking in a dynamic physical environment. The poor throughput is a direct consequence of the von Neumann architecture bottleneck and the massive overhead introduced by the operating system's task scheduler. The CPU is forced to execute the CNN matrix multiplications sequentially, continuously fetching instructions, decoding them, and reading individual RSSI floating-point values from cache hierarchies. This sequential execution model fundamentally limits the parallel nature of convolutional filtering. This inherent limitation motivated the transition to a dedicated, custom hardware accelerator. By designing a heavily pipelined, spatially unrolled Multiply-Accumulate (MAC) core tailored specifically for 2D Radio Map grids, the architecture completely bypasses the sequential instruction bottleneck, trading flexibility for immense operational density to deliver real-time, high-throughput localization predictions.

2. Roofline Analysis

Understanding the theoretical performance limits of the target hardware required a rigorous roofline analysis. The mathematical profile of the system is dictated by its arithmetic intensity (AI), which measures the ratio of floating-point (or integer) operations performed to the total bytes of memory traffic required to sustain those operations. The theoretical AI of the target CNN kernel for our indoor localization model was calculated to be exactly 1.058 FLOP/byte.

In Milestone 1, our initial theoretical performance target was set at an ambitious 80 GFLOP/s, assuming a heavily parallelized systolic array operating at high frequencies. However, plotting the roofline model quickly revealed a severe memory wall. Sustaining 80 GFLOP/s at a low AI of 1.058 FLOP/byte requires a massive, sustained external memory bandwidth of 75.61 GB/s ($80 \text{ GFLOP/s} \div 1.058 \text{ FLOP/byte} \approx 75.61 \text{ GB/s}$).

When mapping this operational requirement onto our target FPGA System-on-Chip (SoC) host platform—which utilizes standard AXI4-Stream interfaces to bridge the processing system (PS) and programmable logic (PL)—the bottleneck shifted dramatically from compute limitations to memory limitations. A standard, wide 512-bit AXI4-Stream interface clocked at a rapid 300 MHz provides a theoretical peak bandwidth of only 19.2 GB/s. Because the mathematical requirement of 75.61 GB/s vastly exceeds the hardware's 19.2 GB/s maximum physical capacity, the design is severely interface-bound. Without targeted architectural intervention, the roofline model dictates that the attainable performance of the kernel would be hard-capped by the memory interface at a mere 20.31 GFLOP/s ($19.2 \text{ GB/s} \times 1.058 \text{ FLOP/byte} \approx 20.31 \text{ GFLOP/s}$).

This severe memory bandwidth bottleneck directly shaped the final datapath and memory hierarchy of the architecture. To circumvent the AXI4-Stream limitations, the hardware could not simply request new data for every calculation. Instead, the architecture was forced to prioritize deep on-chip data reuse. By leveraging local SRAM caching mechanisms to store the RSSI grids internally, the effective arithmetic intensity is artificially inflated, decoupling the high-speed compute arrays from the sluggish external bus limits. (See Figure 1: Final Roofline Model for visual confirmation of this performance bounding).




Figure 1: Final Roofline Model showing the gap between the software baseline and hardware accelerator.

3. Precision and Data Format

To further alleviate the intense memory bandwidth bottleneck identified in the roofline analysis and to drastically minimize the standard cell logic footprint of the physical layout, strict INT8 (8-bit signed integer) quantization was adopted for all RSSI feature inputs and pre-trained network weights. In conventional general-purpose computing, machine learning models utilize 32-bit floating-point representations (FP32), which offer an immense dynamic range but require highly complex silicon structures—such as mantissa multipliers, exponent adders, and barrel shifters—to execute. By contrast, an INT8 representation completely drops this overhead.

The transition to INT8 achieves two vital goals. First, it reduces the required external memory bandwidth by a factor of four compared to FP32, allowing the AXI4-Stream bus to transfer four discrete RSSI data points per cycle rather than one. Second, INT8 multipliers require a fraction of the combinational logic gates required by FP32 hardware. This allows four times as many independent MAC operations to be processed per clock cycle within

the exact same physical area and dynamic power budget, which is a critical necessity for maximizing prediction throughput in power-constrained edge deployment environments.

However, multiplying 8-bit values and continuously accumulating them across a highly parallelized 2D grid inherently introduces the risk of arithmetic overflow, which would catastrophically corrupt the localization output. To safely prevent overflow during these extensive parallel dot-product operations, the internal accumulation architecture was decoupled from the 8-bit format and widened to a robust 32-bit signed integer register. As verified comprehensively during the Milestone 2 design phase using our Python baseline model (`test_localization.py`), the quantization noise introduced by clamping the raw analog RSSI values to an 8-bit resolution did not unacceptably degrade the spatial localization accuracy. This mathematical modeling formally proved INT8's viability as the optimal, resource-efficient data format for this accelerator pipeline.

4. Dataflow and Architecture

To effectively and efficiently process the continuous stream of 2D RSSI spatial maps, the accelerator datapath architecture was specifically built around a weight-stationary dataflow pattern. In a weight-stationary architectural model, the ILS neural network weights are pre-loaded into local, dedicated registers immediately adjacent to the arithmetic computation units. These weights are held entirely stationary while the streaming RSSI feature data flows iteratively past them. This specific dataflow pattern is the optimal choice for our specific application kernel because the CNN weights for the indoor environment map remain constant and immutable during any active prediction phase. Conversely, the incoming RSSI sensor data is highly volatile and continuously updated as the target navigates the physical space. By keeping the static weights localized to the multipliers, the architecture completely avoids the necessity of repeatedly fetching the same network coefficients from external RAM. This architectural choice addresses the core vulnerability exposed in the roofline analysis, drastically reducing overall memory bandwidth requirements and allowing the external bus to be entirely dedicated to streaming in fresh sensor data. If an output-stationary model had been chosen instead, the architecture would be forced to constantly stream partial sums back and forth to memory, immediately causing bus saturation and halting the calculation pipeline. (See Figure 2: Algorithm Dataflow Diagram).

CNN-based Indoor Localization Pipeline




Figure 2: Algorithm Dataflow Diagram detailing the CNN spatial processing kernel.

The core compute engine is constructed from a matrix of parallel INT8 Multiply-Accumulate (MAC) arrays. Each individual MAC unit consists of an 8-bit combinational multiplier directly feeding into the 32-bit accumulator register. The internal data path is heavily pipelined, allowing the computation engine to accept a new set of inputs and produce a mathematically updated partial sum on every single rising clock edge. The memory hierarchy is orchestrated by a top-level wrapper module, `localization_core.sv`, which actively manages the local SRAM buffers to cache incoming sensor grids and sequence them appropriately before feeding them into the active MAC array matrix. (See Figure 3: System Block Diagram).




Figure 3: System Block Diagram illustrating the AXI4-Stream integration.

5. Hardware Interface

To ensure seamless integration with the host FPGA System-on-Chip (such as a Xilinx Zynq platform) and to maintain compatibility with surrounding commercial intellectual property (IP) blocks, the accelerator explicitly implements the industry-standard AXI4-Stream interface. The AXI4-Stream protocol was chosen over a standard AXI4 Memory-Mapped interface because its simplified valid/ready handshaking mechanism perfectly supports the continuous, unidirectional flow of streaming pixel data required by our pipelined MAC arrays. By dropping the requirement for explicit read and write address channels, the protocol minimizes signaling overhead and reduces routing congestion in the physical layout. Furthermore, the interface effectively utilizes tlast frame indicators to explicitly demarcate the boundaries of the streaming 2D data grids, allowing the internal control logic to easily track the start and end of individual spatial maps.

As established mathematically during the Milestone 1 performance target assessment, reaching the theoretical peak of 80 GFLOP/s at an arithmetic intensity of 1.058 FLOP/byte demands a staggering 75.61 GB/s of effective bandwidth. Because a standard 512-bit AXI4-Stream interface clocked at 300 MHz caps out at roughly 19.2 GB/s, the design remains fundamentally and undeniably interface-bound. However, by implementing the weight-stationary architecture and deep local caching mechanisms described in the previous section, the internal effective bandwidth of the MAC arrays operates independently of the AXI4-Stream bottleneck. The external interface may be constrained to 19.2 GB/s, but the

internal pipeline remains fully saturated, ensuring that no compute cycles are wasted waiting for the bus.

6. Verification

Ensuring both the exact mathematical correctness and the physical timing integrity of the hardware was managed through a rigid, multi-tiered verification strategy designed to eliminate logic discrepancies before physical synthesis. First, a golden software model was developed (`test_localization.py`). This Python-based algorithmic testbench heavily utilized the NumPy library to simulate the exact INT8 quantization constraints, clamping behavior, and 32-bit accumulations that would physically occur in the silicon hardware. Executing this code generated a mathematical ground truth vector for all subsequent hardware tests.

Following the generation of the golden vectors, the RTL implementation was verified using robust SystemVerilog testbenches compiled and executed via the open-source Icarus Verilog simulator. The hardware test cases were specifically designed to aggressively stress the AXI4-Stream wrapper logic. The testbench autonomously fed synthesized, fixed-seed RSSI signal data into the active core, actively dropping and raising the tready signal to simulate backpressure and stall conditions, ensuring the pipeline did not drop or duplicate data under duress.

To formally prove the functionality of the architecture, raw simulation toggle data was dumped to uncompressed `.vcd` (Value Change Dump) files. Using GTKWave, visual proof of the physical timing waves was extracted. This waveform analysis explicitly demonstrated flawless valid/ready handshaking across the AXI4-Stream bus and verified that the hardware's two's complement arithmetic execution perfectly matched the expected NumPy golden model integers on a strict, cycle-by-cycle basis. (See Figure 4: Final Hardware Waveform).




Figure 4: Final Hardware Waveform showing valid/ready handshaking.

7. Synthesis Results

The verified RTL design was subsequently synthesized and physically routed targeting the open-source Skywater 130nm (Sky130) Process Design Kit (PDK) utilizing the OpenLane 2 automated physical design flow. The toolchain successfully guided the design through Yosys for complex logic synthesis and mapping, followed by OpenROAD for physical floorplanning, standard cell placement, clock tree synthesis (CTS), and detailed wire routing.

According to the post-routing standard cell statistics report (stat.rpt), the fully synthesized and routed macro consumes a total physical layout area of precisely 10932.985600 square micrometers. The dominant contributor to this extensive logic footprint is the 32-bit accumulator matrix and the highly parallelized INT8 multiplication arrays. Standard cell area within this macro is heavily consumed by the wide, deep combinational adder trees required to mathematically sum the parallel partial products together within the span of a single clock cycle.

The OpenROAD dynamic power analysis report (power.rpt) estimates the total power consumption of the active macro to be 4.495985e-04 Watts (approximately 0.45 milliwatts). Dynamic switching power acts as the absolute dominant contributor to this

total. Because the MAC pipeline is continually saturated with new RSSI streaming data on every valid clock cycle, the switching activity (toggling rate) of the standard cells and the extensive clock tree network is exceptionally high compared to the static leakage power, which remains negligible in the 130nm node.

Crucially, the completed layout achieved successful timing closure without requiring architectural logic compromises. As definitively documented in the maximum delay static timing analysis (STA) signoff report (max.rpt), there are strictly zero setup violations at the newly relaxed target clock period of 20.0 ns (50 MHz). The critical path analysis reveals a highly robust positive setup slack of 11.263489 ns. The maximum data arrival time through the deepest combinational logic path registered at just 4.486512 ns, easily clearing the required time boundary of 15.750001 ns. This substantial timing margin provides a highly stable operating point across the typical process corner (nom_ss_100C_1v60), mathematically ensuring reliable data capture between internal pipeline stages and formally validating the architectural decision to back off from the original, overly aggressive frequency targets.

8. Benchmark Comparison

To quantitatively validate the engineering benefit of the custom hardware architecture, the completely synthesized ASIC macro was benchmarked directly against the initial Milestone 1 software baseline. As previously established, the Python-based software model, executed natively on a high-end Intel Core i7-12700H general-purpose CPU, achieved a maximum sustained throughput of only 183.18 RSSI samples per second.

By stark contrast, the newly routed hardware accelerator—clocked reliably at 50 MHz—features a datapath fully pipelined to accept one valid input and execute one complete MAC arithmetic operation every single clock cycle without stalling. This architectural efficiency results in a verified, measured throughput of 50,000,000 RSSI samples per second.

The mathematical speedup provided by the custom silicon is calculated as follows:

- **Speedup** = (M4 Accelerator Throughput) / (M1 Baseline Throughput)
- **Speedup** = 50,000,000 / 183.18 = **272,955x**

This extraordinary, multiple-orders-of-magnitude performance speedup is a direct physical consequence of abandoning the sequential instruction fetch-decode-execute cycle mandated by a general-purpose operating system in favor of a spatially unrolled, application-specific data path that guarantees execution on every clock edge. Furthermore, the energy efficiency gap between the two platforms is massive. The Core i7 CPU operates

with a substantial 45W Thermal Design Power (TDP) budget merely to sustain its sluggish 183.18 samples/sec throughput. In direct comparison, the custom Sky130 standard cell macro operates in the extreme low milliwatt range (0.45 mW) while simultaneously executing 50 million discrete operations per second.

Bridging the Performance Gap:

Despite this staggering performance speedup over the software baseline, there remains a distinct, highly visible gap between the *theoretical* maximum performance target defined in Milestone 1 and the actual *measured* performance of the physically synthesized M4 macro. The initial M1 roofline analysis targeted a peak computation rate of 80 GFLOP/s based on the early assumption of a massive, heavily parallelized systolic array running at 300+ MHz. However, the measured performance of our synthesized, single-pipeline compute core operating cleanly at 50 MHz yields exactly 0.1 GFLOP/s (100 MFLOP/s).

This gap exists fundamentally because physically unrolling the architecture to concurrently support 800+ parallel MAC units at 300 MHz would severely and catastrophically violate the strict area limits and routing congestion tolerances of the Sky130 standard cell library, while simultaneously exceeding the hard limits of the AXI4-Stream interface memory bottleneck. Scaling the physical hardware from 0.1 GFLOP/s to the theoretical 80 GFLOP/s would necessarily require tiling multiple distinct macros across the die and radically upgrading the memory interface to a multi-channel configuration.

9. What Did Not Work

The iterative journey from high-level Python application code to a physically realized, timing-closed silicon layout exposed several severe architectural and toolchain failures that necessitated significant, highly targeted design revisions.

1. Aggressive Clocking and Setup Violations:

Initially, the physical synthesis flow was configured with a highly aggressive clock period target of 12.0 ns (approximately 83 MHz) in the OpenLane config.json configuration file. At this elevated frequency, OpenROAD's Static Timing Analysis (STA) engine immediately flagged massive, systemic setup time violations ([Checker.SetupViolations]) across all tested process corners (max, min, and nom ss_100C_1v60). The core issue was combinatorial logic depth: forcing the hardware to perform a discrete 8-bit multiplication and a subsequent wide 32-bit addition all within the span of a single clock tick fundamentally exceeded the physical electron transit time available in the older Sky130 technology node.

To achieve closure without forcing a fundamental rewrite of the RTL to manually inject intermediate pipeline registers into the math paths, the clock period was deliberately relaxed to 20.0 ns (50 MHz). This represented a calculated tradeoff, sacrificing peak operating frequency to successfully yield a clean layout. If I were to redesign the core from scratch in a future revision, I would insert an explicit hardware pipeline register squarely between the multiplication and accumulation logic clouds to break the critical path, allowing the design to easily close at 100+ MHz.

2. Toolchain and Environment Instability:

Navigating the OpenLane 2 open-source automation ecosystem presented distinct operational differences and hurdles compared to working within commercial physical design environments such as Cadence Innovus or Synopsys Genus. The local Linux subsystem environment frequently lost the executable path to the core OpenLane binaries, continuously throwing command not found terminal errors during repeated synthesis attempts. It was ultimately discovered that the OpenLane structural framework was deeply and inextricably tied to an isolated Nix package manager environment. Attempting to force execution via absolute binary paths or standard Python venv activations routinely failed. The permanent solution required explicitly invoking the nix-shell command directly inside the tool's root installation directory to initialize the hidden paths before navigating back to the project workspace to run the flow logic.

3. OpenROAD Constraining Quirks:

The automated flow generated several pervasive warnings regarding undefined SDC (Synopsys Design Constraints) files and missing physical power pad locations (VSRC_LOC_FILES). Because manually customized .sdc files were not provided to the toolchain, OpenROAD was forced to fall back on highly generic, unoptimized fallback constraints, which highly likely exacerbated the initial timing and routing failures. Furthermore, obscure LEF58 enclosure warnings frequently appeared during the detailed routing stage. While these enclosure warnings did not fatally break the final GDSII geometry generation, they indicate sub-optimal geometries. In a future tape-out revision, I would author a highly dedicated signoff.sdc constraint file to establish explicitly defined input and output delay constraints, multi-cycle constraints, and false paths. Supplying this explicit mathematical data would afford the detailed router significantly better heuristics to tightly optimize the physical standard cell placement and minimize unnecessary wireline propagation delays.